

stock price prediction

evaluating different models

we evaluate and visualize different models to determine the correlation (aka similarity) between their outputs, in order to create diverse strategies that will perform well when combined into one model.

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import yfinance as yf
from sklearn.linear_model import Ridge, Lasso
from sklearn.neighbors import KNeighborsRegressor
from sklearn.model_selection import train_test_split
from sklearn.model_selection import TimeSeriesSplit, cross_val_score
from sklearn.ensemble import VotingRegressor
from sklearn.metrics import make_scorer
```

```
In [ ]: import warnings
warnings.filterwarnings("ignore")
```

creating features

- The spread feature measures volatility.
- The gap feature measures the price movement after hours
- The intraday feature measures the price movement from the open to close.

```
In [ ]: def create_features(df):
    df['Spread'] = df['High'] - df['Low']
    df['Gap'] = df['Open'] - df['Close'].shift(1)
    df['Intraday'] = df['Open'] - df['Close']
    return df
```

calculating features

- to measure the percentage change and rolling average for Spread, Gap, and Intraday over the last n periods.
- lookback: maximum period over which we want to look back. step: interval at which we calculate the new features.

```
In [ ]: def process_features(df, lookback, step):
    for i in range(step, lookback+1, step):
        df['%d Spread' % (i)] = df['Spread'].pct_change( periods=i, fill_method='ffill')
        df['%d Rolling Avg Spread' % (i)] = df['Spread'].rolling(window=i).mean()

        df['%d Gap' % (i)] = df['Gap'].pct_change( periods=i, fill_method=None)
        df['%d Rolling Avg Gap' % (i)] = df['Gap'].rolling(window=i).mean()

        df['%d Intraday' % (i)] = df['Intraday'].pct_change( periods=i, fill_method='ffill')
        df['%d Rolling Avg Intraday' % (i)] = df['Intraday'].rolling(window=i).mean()
    return df
```

dropping extra columns

```
In [ ]: def drop_features(df):
    df.drop(columns=['Spread', 'Gap', 'Intraday'], inplace=True)
    df.drop(columns=['Open', 'High', 'Low', 'Close', 'Volume', 'Adj Close'], inplace=True)
    return df
```

creating a target

- calculates the percentage change in price over these days and then takes the natural logarithm of that change.
- This can help predict future price movements based on past data.

```
In [ ]: def create_target(df, lookforward=2, target='Open'):
    df['Target'] = np.log(df[target].shift( periods=-lookforward)/df[target]).shift(lookforward)
    return df
```

developing models

- Ridge Regression: It smooths out our predictions by putting a limit on how much the factors can influence the outcome.
- Lasso Regression: It helps us pick out the most important factors for prediction and ignores the less important ones
- KNeighborsRegressor: A non-linear regression model that makes predictions based on the closest data points.

```
In [ ]: estimator1 = Ridge()
estimator2 = Lasso(alpha=.001)
estimator3 = KNeighborsRegressor()
models = [estimator1, estimator2, estimator3]
```

```
In [ ]: lookforward = 2  
        step = 21  
        lookback = 21  
        results = pd.DataFrame()
```

defining the models

- SPY (representing the S&P 500) reflects the stock market, while AGG (a bond market index) reflects the bond market.
- Combining them provides insights from both markets, which can enhance the predictive power of the model.

```

In [ ]: for model in models:
    if lookback >= step:
        # importing data
        spy = yf.download('SPY', start='2004-01-01')
        agg = yf.download('AGG', start='2004-01-01')

        # creating target
        spy = create_target(spy, lookforward, target='Open')

        # adding features
        spy = create_features(spy)
        spy = process_features(spy, lookback, step)
        spy = spy.add_suffix(' SPY')

        agg = create_features(agg)
        agg = process_features(agg, lookback, step)
        agg = agg.add_suffix(' AGG')
        cv = pd.merge(spy, agg, how='inner', on='Date')

        # remove rows from end and begining to enusre enough data for lookback
        cv.drop(cv.tail(lookforward).index, inplace=True)
        cv.drop(cv.head(lookback).index, inplace=True)

        # filling null values
        X = cv
        y = X[['Target SPY']]
        X = X.drop(columns=['Target SPY'])
        X.fillna(method="ffill", inplace=True)
        X.replace([np.inf, -np.inf], 0, inplace=True)
        X.fillna(0, inplace=True)

        #training the model
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.

        model.fit(X_train, y_train)
        y_pred = model.predict(X_test)
        results[model] = y_pred.flatten()

    else:
        print('Lookback must be greater than or equal to step')
        break

```

visulaise correltation of models

```

In [ ]: !pip install seaborn --upgrade -q

```

```

In [ ]: correlation = results.corr()
sns.heatmap(correlation, annot=True, fmt='.2f', cmap='coolwarm', center=0)
plt.tight_layout()
plt.show()

```

- Here we see the models are mostly uncorrelated, which is good.

- when we have uncorrelated models, each model has its own strategy to maximize the scoring function, and when we combine all these uncorrelated strategies, we get a meta-model which has a robust strategy

Finding optimal interval and lookback

create a scoring function

we are scoring different models on their average returns,

```
In [ ]: def custom_score(y_true, y_pred):
        pred_sign = np.sign(y_pred)
        y_true = np.squeeze(y_true)
        returns = np.where((pred_sign == 1), y_true, 0)
        return returns.mean()

        custom_scorer = make_scorer(custom_score, greater_is_better=True)
```

combining 3 models into 1

We combine the 3 models into 1 using the VotingRegressor. The VotingRegressor works by averaging the output of the models that its composed of into a single output.

```
In [ ]: estimator1 = Ridge()
        estimator2 = Lasso(alpha=.001)
        estimator3 = KNeighborsRegressor()
        models = [estimator1, estimator2, estimator3]
        estimator = VotingRegressor(estimators=[('Ridge', estimator1), ('Lasso', estimator2), ('KNeighborsRegressor', estimator3)])
```

Define target, cross validation folds, interval, and lookback parameters.

- we test the model 5 times (n_splits=5). We also have to add a gap between each split, otherwise there will be some overlap, which could cause the models to peek at data it shouldn't, and artificially inflate its score.
- Next we define the steps (or intervals) and lookbacks we want to test our model at, in order to find the ones that it scores the highest on.
- we create an object to store the results of our model in.

```
In [ ]: lookforward = 2
        tscv = TimeSeriesSplit(n_splits=5, gap=lookforward)
        steps = [1, 2, 3, 4, 5]
        lookbacks = [1, 2, 3, 4, 5]
        cv_results = pd.DataFrame(columns=['step', 'lookback', 'score'])
```

Evaluate and visualize scores of different combinations of parameters

```

In [ ]: for step in steps:
        for lookback in lookbacks:
            if lookback >= step:
                # importing data
                spy = yf.download('SPY', start='2004-01-01')
                agg = yf.download('AGG', start='2004-01-01')

                # creating target
                spy = create_target(spy, lookforward, target='Open')

                # adding features
                spy = create_features(spy)
                spy = process_features(spy, lookback, step)
                spy = spy.add_suffix(' SPY')

                agg = create_features(agg)
                agg = process_features(agg, lookback, step)
                agg = agg.add_suffix(' AGG')
                cv = pd.merge(spy, agg, how='inner', on='Date')

                # remove rows from end and beginning to ensure enough data for look
                cv.drop(cv.tail(lookforward).index, inplace=True)
                cv.drop(cv.head(lookback).index, inplace=True)

                # filling null values
                X = cv
                y = X[['Target SPY']]
                X = X.drop(columns=['Target SPY'])
                X.fillna(method='ffill', inplace=True)
                X.replace([np.inf, -np.inf], 0, inplace=True)
                X.fillna(0, inplace=True)

                cv_score = cross_val_score(estimator=estimator, X=X, y=y,
                                           scoring=custom_scorer,
                                           cv=tscv, verbose=0)

                new_row = {'step': step, 'lookback': lookback, 'score': cv_score.mean()}
                cv_results = pd.concat([cv_results, pd.DataFrame([new_row])], ignore_index=True)

```

visualise heatmap

```

In [ ]: cv_results = cv_results.pivot(index='step', columns='lookback', values='score')
sns.heatmap(cv_results, annot=True, fmt='.6f', linewidths=.25, cmap='Blues')
plt.title('Net Average Return')
plt.show()

```

Here we can see a 1 period step interval with a maximum look back of 2 periods performs the best. In the next notebook, we will adjust the parameters of the models

In []:

In []:

```
In [ ]: !pip install backtesting -q
        from backtesting import Strategy, Backtest
```

```
In [ ]: step = 1
        lookback = 2

        estimator1 = Ridge(alpha=0.002840026017965097)
        estimator2 = Lasso(alpha=0.002999525583333498)
        estimator3 = KNeighborsRegressor(n_neighbors=8)
        models = [estimator1,
                   estimator2,
                   estimator3,]
        estimator = VotingRegressor(estimators=[('Ridge', estimator1),
                                                ('Lasso', estimator2),
                                                ('KNN', estimator3)],)

        X_test = X_test.iloc[(abs(lookforward)):]
        y_test = y_test.iloc[(abs(lookforward)):]

        estimator.fit(X_train, y_train)
        forecasted = estimator.predict(X_test)

        data = yf.download('SPY', start='2004-01-01')
        data.drop(data.tail(lookforward).index,inplace=True)
        data.drop(data.head(lookback).index,inplace=True)
        data = data.iloc[(-X_test.shape[0]):]
        data['forecastedValue'] = forecasted
        prediction = data
```

```
In [ ]: class MyStrategy(Strategy):
        Data = prediction

        def init(self):
            super().init()

        def next(self):
            if self.data.forecastedValue < 0:
                self.sell()
            elif self.data.forecastedValue > 0:
                self.buy()

bt = Backtest(prediction, MyStrategy,
              cash=1000,
              trade_on_close=False,
              exclusive_orders=True
              )
print(bt.run())
```

```
In [ ]:
```